

System Design Day 2: Database Design

SQL (Structured Query Language) vs NoSQL (Not Only SQL) with Sharding

◆ 1. Data Structure

- **SQL (e.g., MySQL):**
 - Structured, relational data organized into **tables** with rows and columns.
 - Relationships are enforced using **foreign keys**.
 - **NoSQL (e.g., DynamoDB):**
 - Flexible data models like **key-value**, **document**, or **wide-column**.
 - Best for loosely structured or rapidly changing data.
-

◆ 2. Schema

- **SQL:**
 - Enforces a **fixed schema**; changes need explicit migration.
 - Promotes **data integrity and validation**.
 - **NoSQL:**
 - **Schema-less** or flexible schema — fields can vary across records.
 - Supports rapid iteration and deployment.
-

◆ 3. Scalability

- **SQL:**
 - Scales **vertically** (adding more CPU/RAM to a single server).
 - Challenging to scale horizontally without complex partitioning or sharding.
- **NoSQL:**

- Designed for **horizontal scaling** — data is spread across many nodes automatically.
-

◆ 4. Consistency

- **SQL:**
 - Ensures strong consistency with **ACID** properties:
(*Atomicity, Consistency, Isolation, Durability*)
 - **NoSQL:**
 - Uses **BASE** properties:
(*Basically Available, Soft state, Eventual consistency*)
 - Consistency can be relaxed to improve availability and performance.
-

◆ 5. Query Capabilities

- **SQL:**
 - Rich **query language (SQL)**: supports joins, aggregates, filtering, subqueries.
 - Ideal for complex reporting and data relationships.
 - **NoSQL:**
 - Limited to **simple, key-based** access.
 - Optimized for performance over flexibility.
-

◆ 6. Use Cases

- **SQL:**
 - **Banking systems**, ERP, E-commerce order management — where consistency and relations matter.
 - **NoSQL:**
 - **User activity tracking**, IoT telemetry, session stores — where speed and scale are critical.
-

◆ 7. Integration with Java

- **SQL:**
 - Typically used with **Spring Data JPA (Java Persistence API)**.
 - Enables ORM (**Object-Relational Mapping**) and declarative transaction handling.
 - **NoSQL:**
 - Integrated using **Spring Data DynamoDB** or native **AWS SDK**.
 - Works well with **microservices** for loosely coupled, independent modules.
-

◆ 8. Sharding

📖 Sharding Definition:

Sharding is a database partitioning technique where large datasets are split into smaller, more manageable pieces called "shards", based on a shard key (e.g., user ID, region). Each shard is stored on a different server or node to distribute load and improve scalability.

- **SQL:**
 - Manual sharding using application logic or middleware.
 - Adds complexity in **query routing and joins across shards**.
 - **NoSQL:**
 - **Native support** for automatic sharding.
 - Example: DynamoDB uses **partition key + sort key** for transparent and automatic data distribution.
-

◆ 9. Performance at Scale

- **SQL:**
 - Performance can degrade under **high concurrent load** or large dataset size.
 - Needs caching layers or read replicas.

- **NoSQL:**
 - Designed for **low-latency** and **high-throughput** workloads.
 - Optimized for write-heavy applications.
-

Real-Life Application Examples

SQL Example: E-Commerce Order Management System (MySQL)

- **Why SQL?**
 - Requires **strong consistency** and **complex relationships** (orders → users → payments).
 - Relational queries needed for inventory, reporting, and analytics.
- **Tech Stack:**
 - MySQL + Spring Boot + Spring Data JPA
 - Scaled using **Read Replicas**, **Database Connection Pooling**, and **Caching** with Redis.

Scalable Design Patterns Used in SQL Systems

- **Read Replicas:** Distribute read load from the primary DB.
 - **Database Connection Pooling:** Reduces overhead from opening/closing connections.
 - **CQRS (Command Query Responsibility Segregation):** Split read and write models for performance.
 - **Caching Layer (Redis, Ehcache):** Reduce DB load on frequently accessed data.
-

NoSQL Example: Social Media User Activity Tracking (DynamoDB)

- **Why NoSQL?**
 - Handles massive, fast writes (likes, posts, follows).
 - Schema-less design supports evolving data.
 - Scales horizontally for millions of users.

- **Tech Stack:**
 - DynamoDB + AWS Lambda + Spring Boot (for APIs)
 - Uses **partition key** (user ID) + **sort key** (timestamp).

Scalable Design Patterns Used in NoSQL Systems

- **Auto-Sharding:** Managed by DynamoDB via partition keys.
- **Event Sourcing:** Store each user action as an event for audit or replay.
- **Time-Series Data Modeling:** Store data with timestamps for range queries.
- **Microservices:** Decompose features by bounded context; each owns its own NoSQL data store.

How Sharding Affects Query Performance

◆ 1. Positive Impact on Query Performance (When Done Right)

Benefit	Explanation
Parallel Query Execution	Queries can be executed in parallel across shards , reducing response time.
Reduced Data Size per Node	Each shard holds less data , so index scans and lookups are faster .
Less Contention	By distributing load, read/write contention is minimized on any single node.
Improved Scalability	As the dataset grows, you can add more shards , keeping query performance stable.

◆ 2. Negative Impact on Query Performance (When Done Poorly)

Risk	Impact
Cross-Shard Queries	If the query doesn't include the shard key , the system has to query all shards , increasing latency.
Hotspotting	Uneven data distribution (e.g., all writes go to one shard) leads to overloaded shards , slowing performance.

Risk	Impact
Complex Joins/Aggregations	SQL-style joins and aggregates across shards are difficult or slow (often not supported in NoSQL).
Rebalancing Overhead	Adding/removing shards can require data migration , temporarily affecting performance.

Example:

- Suppose you're sharding a **user activity log** by `user_id`.
- Query like `SELECT * FROM activity WHERE user_id = 'abc123'` will hit **only one shard** → **fast**.
- But a query like `SELECT * FROM activity WHERE action = 'login'` has no shard key → needs **scatter-gather across all shards** → **slow**.

Best Practices to Optimize Query Performance in Sharded Systems

1. Design with Shard Key in Mind

Always **include shard key in queries** wherever possible.

2. Avoid Cross-Shard Joins

Denormalize data if needed; joins across shards are expensive.

3. Uniform Data Distribution

Choose a shard key that results in **balanced load and storage**.

4. Use Routing Layers

Implement smart query routers or service layers that know where to send a query.

5. Monitor and Re-shard Carefully

Periodically evaluate shard balance; re-shard during low traffic windows.