

System Design Dat 6: API Gateway Design

◆ What Is an API Gateway?

An **API Gateway** is a **central entry point** for client requests in a **microservices architecture**. It acts as a reverse proxy that routes requests to backend services and handles cross-cutting concerns like:

- Authentication
 - Rate Limiting
 - Routing
 - Caching
 - Logging
 - TLS Termination
-

Evolution of Architecture

1. 2000s – Monoliths

- Clients talked to a single application.
- Simple and easy to deploy, but not scalable.

2. 2010–2012 – Microservices Era

- Applications were split into many services.
- Clients either needed to know each service's URL or use a proxy service—both messy solutions.

3. 2013–2014 – API Gateway Emerges

- Introduced a thin centralized layer to handle all requests.
 - Clients only need to know **one endpoint**.
 - The gateway routes the request to the right microservice.
-

Why Use an API Gateway?

Instead of duplicating the same code in every microservice, the API Gateway handles **common infrastructure responsibilities**:

Routing

- Routes incoming requests to the correct backend service using path-based or host-based rules.

Java Context: Spring Cloud Gateway provides configurable routing using YAML or Java DSL.

Nginx/Ingress: Nginx or Kubernetes Ingress controllers can act as layer-7 routers.

Authentication & Authorization

- Validates tokens (e.g., OAuth2 JWT) and restricts access based on roles/policies.

Java: Spring Security filters can be applied at the gateway level.

Nginx: Can integrate with external OAuth2 proxy or Lua scripts.

Ingress: Can use annotations and plugins like OIDC for auth.

Rate Limiting

- Protects services from abuse and ensures fair usage.

Java: Spring Cloud Gateway supports Redis-based rate limiting via RequestRateLimiter filter.

Nginx: Supports rate limiting via built-in modules (`limit_req`).

Ingress: Ingress-NGINX supports rate limiting with annotations.

Request/Response Transformation

- Converts protocols (e.g., gRPC → REST) or rewrites headers and paths.

| Useful for interoperability and versioning strategies.

TLS Termination

- Offloads HTTPS decryption from microservices.

| Done at API Gateway, Nginx, or Ingress controller.

Other Features

- TLS Termination
- Logging & Monitoring
- Response Transformation (e.g., gRPC → REST)
- Caching of frequent responses

Core Responsibilities of an API Gateway

1. Request Validation

- Ensures requests have valid structure, headers, tokens.

2. Middleware Execution

- Handles authentication, authorization, rate limiting, logging, etc.

3. Routing to Services

- Uses a configuration to map URL paths to backend services.

4. Response Transformation

- Converts protocol-specific responses (like gRPC) to RESTful JSON if needed.

API Gateway Implementations

Managed Services

- **AWS API Gateway**

- **Azure API Management**
- **Google Cloud API Gateway**

Open Source & Self-Managed

- **Spring Cloud Gateway** (Java-native, ideal for Spring Boot ecosystems)
- **Kong, Tyk, Express Gateway**
- **Nginx** (high-performance reverse proxy with custom configuration)
- **Kubernetes Ingress Controller** (like NGINX Ingress or Istio Gateway)

NGINX and Kubernetes Ingress Controllers

NGINX

- Acts as a high-performance reverse proxy and can function as an API gateway.
- Supports:
 - Rate limiting
 - SSL termination
 - Header rewriting
 - Load balancing
 - Auth integration via Lua scripts or external tools

```
nginx
CopyEdit
location /api/ {
    proxy_pass http://backend_service;
    limit_req zone=api_zone burst=10 nodelay;
    auth_request /auth;
}
```

Kubernetes Ingress + NGINX

- Ingress is the K8s-native way to define API gateway rules.

- Works with Ingress Controllers like **NGINX**, **Istio**, or **Traefik**.
- Allows declarative config via annotations:

```
yaml
CopyEdit
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
  annotations:
    nginx.ingress.kubernetes.io/limit-connections: "1"
    nginx.ingress.kubernetes.io/auth-url: "https://auth.myapp.com"
spec:
  rules:
    - http:
        paths:
          - path: /api/users
            pathType: Prefix
            backend:
              service:
                name: user-service
                port:
                  number: 80
```

System Design Interview Tip

| In system design interviews:

- Include an **API Gateway** by default.
- Mention it handles:
 - Routing
 - Auth
 - Rate limiting
 - Response transformation

- Say **"this is implemented via Spring Cloud Gateway or NGINX/Ingress"**, and **move on quickly**.

Spending too much time on the gateway is discouraged—it's considered infrastructure hygiene.



Java + Spring Example

```
yaml
CopyEdit
# application.yml for Spring Cloud Gateway
spring:
  cloud:
    gateway:
      default-filters:
        - AddRequestHeader=X-Request-Source, Gateway
        - RequestRateLimiter=replenishRate=10,burstCapacity=20
      routes:
        - id: user-service
          uri: http://localhost:8081
          predicates:
            - Path=/users/**
          filters:
            - JwtAuthenticationFilter
```

With Spring Security:

```
java
CopyEdit
@EnableWebSecurity
public class SecurityConfig extends WebSecurityConfigurerAdapter {
    @Override
    protected void configure(HttpSecurity http) throws Exception {
        http
            .authorizeRequests()
            .antMatchers("/users/**").hasRole("USER")
            .anyRequest().authenticated()
    }
}
```

```
.and()  
.oauth2ResourceServer()  
.jwt();  
}  
}
```