

# System Design Day 7: Distributed Caching Basics

## ◆ What Is Caching?

**Caching** is the process of storing frequently accessed data in temporary memory to reduce the time and resources required to access it from the original data source (like a database or an API).

---

## Centralized Caching (Traditional Approach)

### Overview:

- Uses a **single cache server** to store frequently accessed data.
- Application checks the central cache before querying the main database.

### Pros:

- Simple setup and easy to manage.
- Reduces latency for repeat data fetches.

### Cons:

- **Single point of failure.**
  - **Scalability bottleneck** – can't handle high concurrent load.
  - Not resilient to node or network failures.
- 

## Distributed Caching (Modern Approach)

### What It Is:

A **distributed cache** stores data across multiple nodes in a network using a **Distributed Hash Table (DHT)**. Nodes are horizontally scalable and automatically rebalance data and traffic.

## ✅ Advantages:

- High availability
- Fault tolerance
- Horizontal scalability
- Low latency and fast response

## ⚙️ Key Features

### ✅ Consistency

- Ensures data accuracy across distributed nodes.
- Some tools (e.g., **Redis Cluster**) provide **eventual consistency**, while others like **Apache Ignite** offer **strong consistency** via transactional and ACID-compliant caches.

### 🔄 Eviction Policies

- Automatically remove stale or least-used items to free memory.
- Common strategies:
  - **LRU (Least Recently Used)**
  - **LFU (Least Frequently Used)**
  - **TTL (Time-To-Live)**

Most distributed caches support configuring eviction based on time, size, or usage.

## 🧱 Multi-Node Setup

- Nodes collaborate and replicate data.
- Load is distributed evenly.
- In case of node failure, others take over (high availability).
- New nodes auto-join and data is rebalanced dynamically.

---

## 🔧 Popular Distributed Caching Tools

### 💠 Redis

- In-memory key-value store.
- Supports data structures (Lists, Sets, Hashes).
- Offers pub/sub and basic persistence.
- Great for Spring Boot apps using `spring-boot-starter-data-redis`.

### ◆ **Memcached**

- Lightweight key-value store.
- High performance, stateless, distributed.
- No advanced structures or persistence.

### ◆ **Hazelcast**

- In-memory data grid.
- Java-native.
- Supports distributed maps, queues, and executors.
- Easily integrated with Spring.

### ◆ **Apache Ignite** ✅

Yes, Apache Ignite is a powerful distributed caching solution — and much more.

#### 🔧 **Features:**

- **In-memory distributed cache** (with optional persistence).
- **Strong consistency** with support for **ACID transactions**.
- **SQL support** (run SQL queries over distributed cache).
- **Compute grid & distributed database** capabilities.
- **Automatic rebalancing** and **partitioning** of data across nodes.
- **Supports near-cache** for local fast access in co-located applications.

#### 🔌 **Spring Integration:**

- Fully compatible with **Spring Cache abstraction** ( `@Cacheable` , `@CacheEvict` ).
- Can be embedded or used in a client/server mode.

- Example:

```
@Cacheable(value = "products", key = "#productId")
public Product getProductById(String productId) {
    return productRepository.find(productId);
}
```

```
spring:
  cache:
    type: ignite
ignite:
  config: classpath:ignite-config.xml
```

## 💡 Use Cases for Distributed Caching

- **Session management** (user tokens, preferences)
- **Reduce database load** with frequently queried data
- **Accelerate microservice communication** (avoid service-to-service calls)
- **Real-time analytics / leaderboards / notifications**
- **Storing reference data** (product info, currency rates, etc.)

## 🧠 Design Considerations

Concern	Design Consideration
<b>Consistency</b>	Choose between <b>eventual</b> , <b>strong</b> , or <b>transactional</b> consistency.
<b>Eviction</b>	Configure TTLs, LRU/LFU, or size limits to avoid memory overflow.
<b>Fault Tolerance</b>	Ensure cluster replication and automatic failover.
<b>Persistence</b>	Use in-memory only (for speed) or with backup (for durability).
<b>Sharding</b>	Tools like Redis Cluster, Ignite, and Hazelcast handle it natively.

## 💻 Java & Spring Boot Integration

Use **Spring Cache abstraction** to decouple business logic from caching implementation. Backends can be swapped (Redis, Hazelcast, Ignite, etc.)

without changing code.

### Example:

```
@Cacheable("users")
public User getUser(String userId) {
    return userRepo.findById(userId);
}
```

### Backends:

- `spring-boot-starter-data-redis`
  - `spring-boot-starter-cache` + Apache Ignite or Hazelcast
- 

## Final Takeaways

- **Centralized caching** is simple but doesn't scale.
- **Distributed caching** is the industry standard for **high availability**, **horizontal scalability**, and **resilience**.
- **Apache Ignite** is a feature-rich caching + compute platform with **strong consistency** and **ACID transactions**.
- Use **Java Spring Cache abstraction** to plug into any caching backend.
- Choose tools and configs based on your **use case**, **consistency requirements**, and **performance needs**.