

System Design Day 11: Ride-Sharing system

Ride-Sharing System Design Summary

Core Functional Requirements

- Rider selects a pickup location via map.
 - Sees ETA and price.
 - Can pay for the ride.
 - Gets matched to a driver.
 - Receives real-time driver updates.
-

High-Level Architecture Components

1. **Clients:** Riders and drivers (mobile/web).
2. **APIs:** Rider & Driver APIs behind Load Balancers; support horizontal scaling.
3. **Databases:**
 - **User DB:** Info about riders & drivers.
 - **Trips DB:** Trip status, route, fare, etc.
 - Use **sharding** on `user_id` and `trip_id` for scalability.
 - Use **global indexes** (by driver_id/rider_id) for fast queries.
4. **Event Bus (Kafka):**
 - Enables decoupling and async communication.
 - Used for real-time data updates (e.g., driver location, ride requests).
5. **Map Service (Mapbox/Google Maps):**
 - Serves map tiles, geocoding, directions.
6. **Payment Integration (e.g., Stripe):**

- Stripe processes payment and invokes a webhook.
- Webhook pushes event to Kafka for ride updates.

7. Pricing Service:

- Implements **dynamic pricing** (surge based on demand).
- Uses **Spark Structured Streaming** on Kafka events → Redis cache.

8. Matching Service:

- Matches riders with nearby drivers.
- Uses **Uber's H3 hexagonal geospatial index** for location clustering.

Real-Time Data Flow

- Use **WebSockets** (vs. polling) for location updates.
- Driver location pushed to rider in real time.
- Kafka event bus distributes updates across services.

Key Design Considerations

- **Scalability:**
 - Load-balanced APIs.
 - Horizontally sharded DB.
- **Fault Tolerance:**
 - Event-driven system allows for async retries.
- **Performance:**
 - Redis cache for pricing and demand.
 - H3 for fast location-based lookups.
- **Extensibility:**
 - Plug-and-play with external services like Stripe and Mapbox.

Example: Ride Lifecycle

1. Rider opens app → sees map via Mapbox.
2. App queries ETA via Ride Service.

3. Price fetched from Pricing API (uses Redis + Spark).
 4. Rider books and pays via Stripe → webhook fires → event added to Kafka.
 5. Matching Service finds nearby driver using H3 cells.
 6. Driver accepts → both parties notified → real-time location sent over WebSocket.
-

Optimization Topics (Not Fully Covered)

- Redundancy and failure recovery.
 - ETA prediction algorithms.
 - Geo-shard placement strategies.
 - Data analytics and BI pipelines.
-

This system showcases **event-driven microservices**, **real-time communication**, **distributed data processing**, and **3rd-party service integration** — making it scalable, modular, and robust for production.